

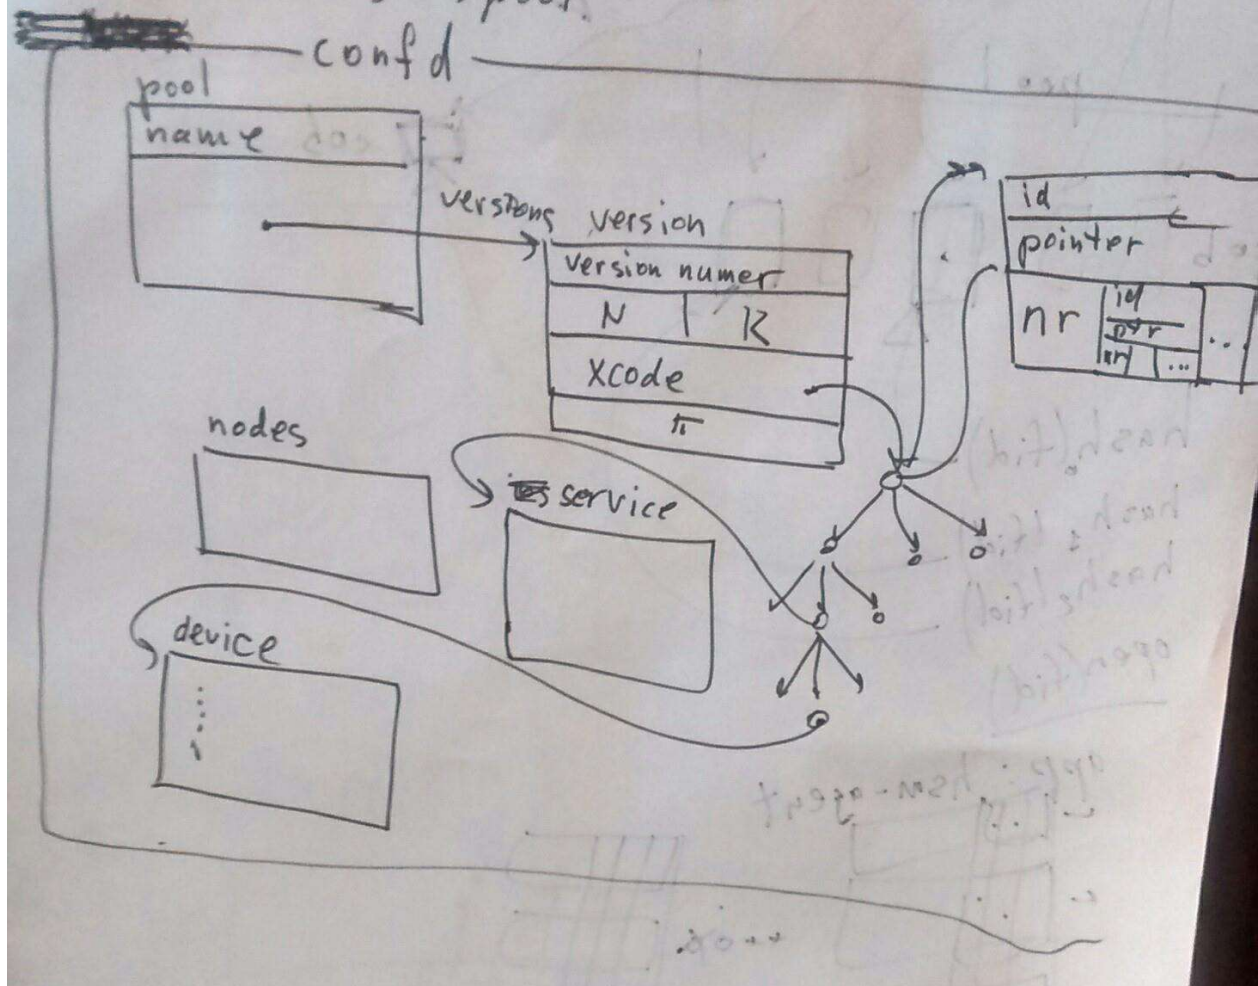
Description

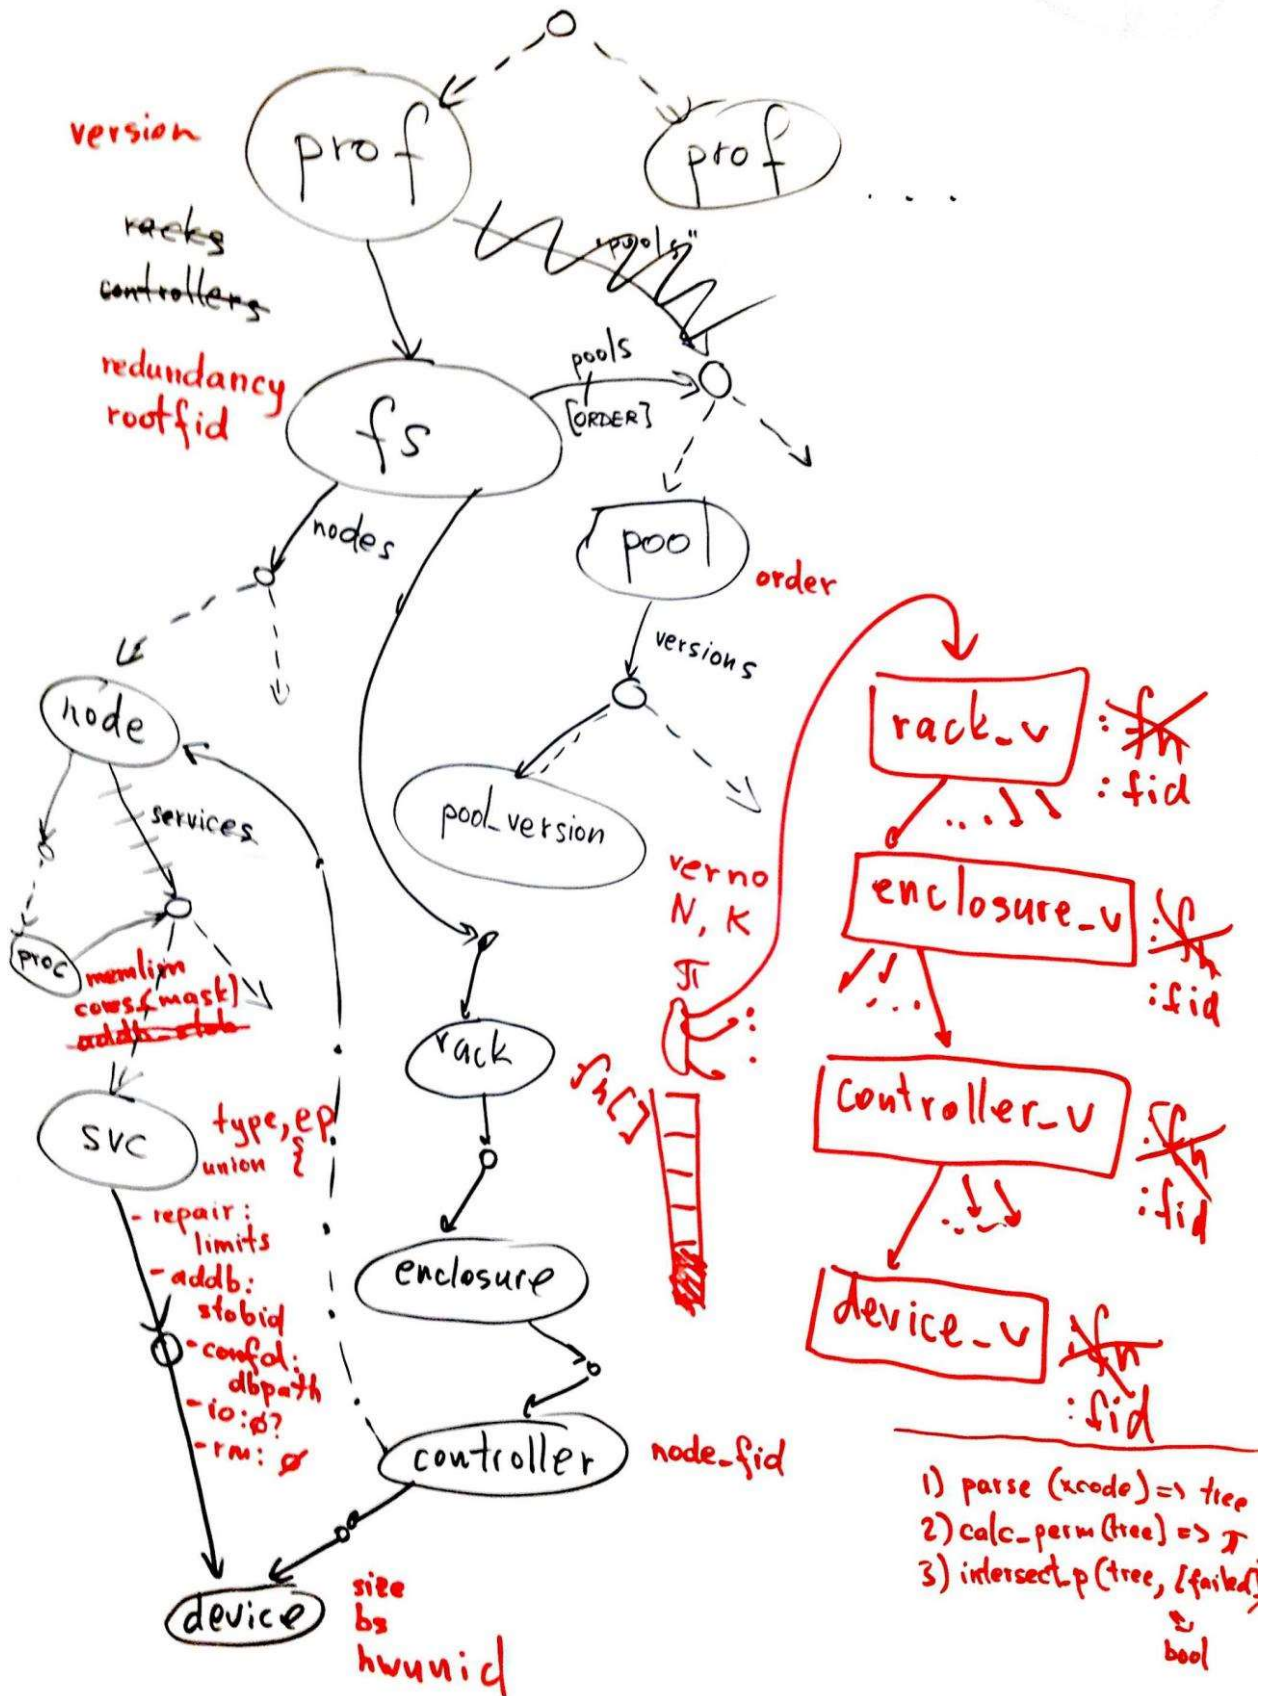
Support multiple pool configurations in order to support dynamic addition of devices, nodes, services, and racks.

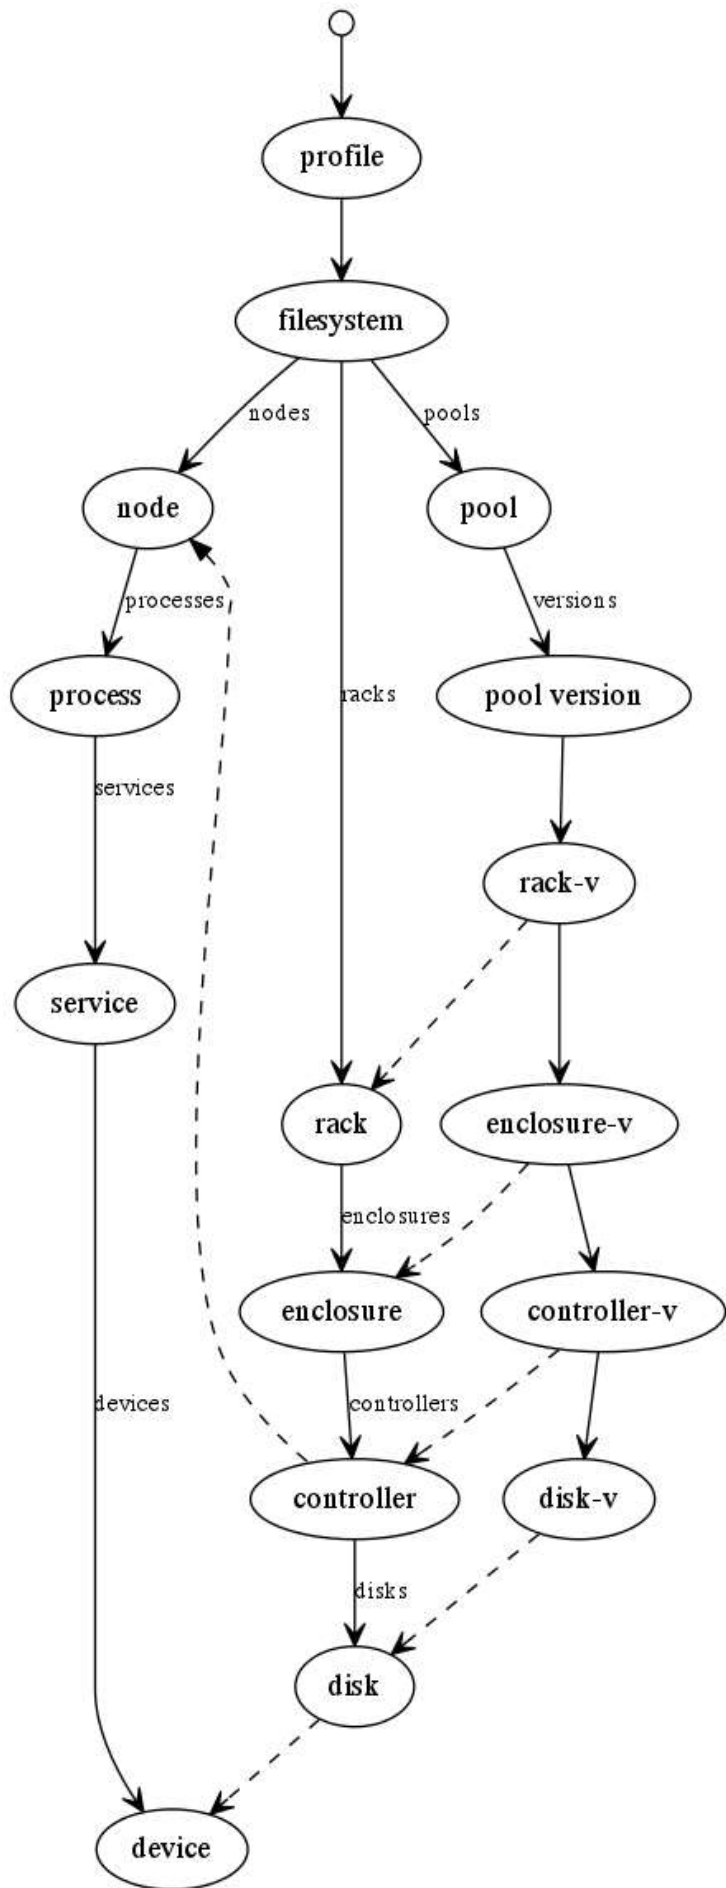
Details

Pools are used to partition hardware resources (devices, servers). Pool versions are used to track changes in pool membership.

1. pool name \rightarrow pool
2. $N+K \rightarrow$ version \rightarrow A version
3. tree of elements (racks, enclosures, controllers, device)
4. list of versions \rightarrow pool







[\[Google Drive folder with DOT source and images\]](#)

Subtasks, estimations

1. Detailed design (DLD) - Mero-531 - (200 loc)
2. add pools to the configuration schema. - (100 loc) MERO-462
3. add pool versions to the configuration schema:
 - a. ~~(not needed, user supplies pool versions) an algorithm to generate pool versions from a pool (if there are C controllers we have 2^C pool versions) (100) -~~
 - b. modify m0t1fs to use multiple pools (100 LOC) - Mero-463
 - c. store the tree associated with a pool version (50). - MERO-464
 - d. modules to parse and manipulate the collection of pool-versions (stored as list) (100). MERO-465
4. Interfaces for m0t1fs to use pools and pool versions - MERO-466
 - a. initialize/update client's failure set (CODE / 50 loc)
 - b. find latest pool/pool version which does not intersect with failure set on configuration changes. (CODE / 50 loc)
5. modify ioservice to work with multiple pools, return version mismatch to clients as necessary, return pool version and pool in getattr reply (CODE / 50 loc) MERO - 467
6. associate layouts with pool versions (CODE / 50) - MERO-468
7. support assignment of pools to newly created objects (CODE / 50 loc) - MERO-469

Dependencies

Finding poolversion:

(Subtask 4b)

```
struct m0_confx_rack{
    struct m0_confx_headerxr_header;
    /* Enclosures on this tack */
    struct arr_fid    xr_encls;
    /** List of associated pool version fids */
    struct arr_fid    xr_pvers;
} M0_XCA_RECORD;
```

```
struct m0_confx_encl{
    struct m0_confx_headerxe_header;
    /* Controllers in this enclosure */
    struct arr_fid    xe_cntrl;
```

```

    /** List of associated pool version fids */
    struct arr_fid    xc_pvers;
} MO_XCA_RECORD;

struct m0_confx_cntrl{
    struct m0_confx_header xc_header;
    /** Storage devices attached with this controller */
    struct arr_fid    xc_devs;
    /** List of associated pool version fids */
    struct arr_fid    xc_pvers;
} MO_XCA_RECORD;

```

pool_version_add() populates
m0_confx_rack::xr_pvers, m0_confx_encl::xe_pvers & m0_confx_cntrl::xc_pvers

Use cases can be modified as :

```

rackv_add(fid, pool_version, rack_fid) {
    rack_v = confc.alloc_add(fid);
    rack_v.rack_fid = rack_fid;
    pool_version.rackv_add(rack_v);

    rack = confc.lookup(rack_fid);
    rack.pvers_add(pool_version);
}

enclosurev_add(fid, pool_version, rack_v, encl_fid) {
    encl_v = confc.alloc_add(fid);
    encl_v.encl_fid = encl_fid;
    rack_v.enclv_add(encl_v);

    encl = confc.lookup(encl_fid);
    encl.pvers_add(pool_version);
}

controllerv_add(fid, pool_version, encl_v, cntlr_fid) {
    cntlr_v = confc.alloc_add(fid);
    cntlr_v.cntlr_fid = cntlr_fid;
    encl_v.cntlr_add(cntlr_v);

    cntl = confc.lookup(cntlr_fid);
    cntl.pvers_add(pool_version);
}

```

```
}
```

```
pool_version_get(confc, failure_set, **pver){
    available_pool_versions = pool_versions;
    for each f_dev in failure_devices{
        /* exclude failed device pool versions */
        available_pool_versions = available_pool_versions - fdev.pvers;
    }
    return available_pool_versions[0];
}
```

```
pool_version_get_nikita(confc, failure_set) : pool_version {
    for (pool in confc.pools) { /* iterate pools in order */
        pool_version = pool.last_version;
        for (e : failure_set) {
            /* e can be device, or rack, or ... */
            if (e.pvers_contains(pool_version)) {
                /* pool_version intersects with the
                failure set, ignore it. */
                continue outer loop;
            }
        }
        return pool_version;
    }
}
```